
Qualitative Probabilistic Programming

Anonymous Author(s)

Affiliation

Address

email

Abstract

When writing probabilistic programs, sometimes it is difficult for the programmer to specify the correct parameterized family of distributions. We explore an extension to probabilistic programming languages that allows programmers to mark some distributions as unspecified. Then, for each distribution, the system can determine some reasonable family of distributions and infer maximum likelihood parameters.

1 Introduction

By separating model specification and inference, probabilistic programming has made it easier for non-experts to implement and use probabilistic models. Practitioners frequently have strong intuitions about the *structure* of their domain knowledge, such as which latent variables exist and what their causal relations are, and probabilistic programming allows them to encode this knowledge. However, it also requires them to specify the specific parametric shape and parameterization of any distributions used, and intuitions tend to be much less precise there. We present Quipp, a system that does *not* require such specification; instead, random variables and random functions can be left undefined and will automatically be filled in under maximum entropy assumptions based on their types and available datasets.

Our formalism can concisely express a wide variety of models that machine learning practitioners care about, and we provide an expectation maximization algorithm that can learn the parameters for many of these models with reasonable efficiency. This system makes it easy for non-experts to encode their beliefs about the data and to get predictions based on as few additional assumptions as possible.

In the following, we first specify the syntax used to write Quipp programs, including the notation for unknown variables and functions. We describe the class of exponential family variables and functions that our system can learn, and present the expectation maximization algorithm used to learn them. We then demonstrate the expressiveness of our language, and the broad applicability of our algorithm, by writing some of the most common machine learning models in Quipp: clustering, naive Bayes, factor analysis, a Hidden Markov model and Latent Dirichlet Allocation.

2 Syntax

Quipp is implemented as a library for webppl programs. Webppl [2] is a probabilistic programming language that is similar to Javascript but also contains features for generating random values, conditioning on values, and estimating expectations. Quipp programs are written as webppl programs that have access to additional special functions.

Here is an example of a Quipp program to cluster 2d points into 2 clusters:

```
var Cluster = Categorical(2);
```

```

054     var Point = Vector(2, Double);
055     var getPoint = randFunction(Cluster, Point);
056
057     var model = function() {
058         var cluster = randomValue(Cluster);
059         observe(getPoint, cluster);
060     };
061

```

We declared two types (`Point` and `Cluster`) and one random function (`getPoint`). Type annotations are necessary for random functions. The type `Vector(2, Double)` expresses the fact that the points are vectors in \mathbb{R}^2 , and the type `Categorical(2)` expresses the fact that there are 2 possible clusters (so a cluster may be either 0 or 1).

The variable `model` specifies a generative model for a single data point. We use the `randomValue` function to generate a random uniform cluster, and then use `getPoint` to generate the point given the cluster. Since `getPoint` is an unknown random function, we will need to infer its parameters. The `observe` function allows us to observe data; here it says that our observation consist of the result of the call `getPoint(cluster)`.

To demonstrate, let us run this example on the Iris dataset, which consists of 150 points [3]. When we run the program on this data, we infer the parameters to the random function `getPoint`. In this case, `getPoint` is a linear function with Gaussian noise, so it will naturally split the data into 2 clusters with equal variance, as seen in Figure 1.




Figure 1: Clustering the Iris dataset

The first cluster is at (6.3, 5.0) and the second is at (5.0, 1.6). They both have a standard deviation of 0.54 in the x direction and 0.69 in the y direction. We could use these parameters to fill in the generative model, giving us a webppl program:

```

101     var model = function() {
102         var cluster = randomInteger(nclusters);
103         return [gaussian(cluster == 0 ? 6.3 : 5.0, 0.54),
104                 gaussian(cluster == 0 ? 5.0 : 1.6, 0.69)];
105     };
106

```

This model is estimated to assign probability density e^{-393} to the data, yielding a perplexity of $e^{-393/150} = 0.0728$.

Table 1: Types in Quipp

T	$\phi_T(X)$	$\psi_T(X)$	randomValue(T)	Random function class
Double	$\begin{bmatrix} X \\ X^2 \end{bmatrix}$	$[X]$	$\mathcal{N}(0, 1)$	Linear regression
Categorical(n)	$\begin{bmatrix} [X = 1] \\ [X = 2] \\ \dots \\ [X = n - 1] \end{bmatrix}$	$\begin{bmatrix} [X = 1] \\ [X = 2] \\ \dots \\ [X = n - 1] \end{bmatrix}$	uniform	n -class logistic regression
Tuple(T1, T2)	$\begin{bmatrix} \phi_{T1}(X_1) \\ \phi_{T2}(X_2) \end{bmatrix}$	$\begin{bmatrix} \psi_{T1}(X_1) \\ \psi_{T2}(X_2) \end{bmatrix}$	$[\text{randomValue}(T1), \text{randomValue}(T2)]$	Independent regression

3 Family of distributions

In the previous example, `randFunction(Categorical(2), Vector(2, Double))` represented 2 unknown axis-aligned 2-dimensional Gaussian clusters with equal variances. In general, `randFunction` returns a randomized function that is a member of some generalized linear model determined by the desired argument types and return type. The distribution of the function’s return value y is some exponential family whose natural parameters are determined from the arguments x :

$$p_\eta(y|x) = \exp(\eta(x)^T \phi(y) - g(\eta(x)))$$

Here, $\eta(x)$ is the natural parameter, $\phi(y)$ is a vector of y ’s sufficient statistics, and g is the log partition function.

To determine $\eta(x)$, we label some subset of the sufficient statistics of both x and y as *features*. For the gaussian distribution, the sufficient statistics are X and X^2 but the only feature is X . For the categorical distribution `Categorical(n)`, the sufficient statistics and features are both $[X = 1], [X = 2], \dots, [X = n - 1]$. The natural parameters corresponding to non-features are constant, while natural parameters corresponding to features are determined as an affine function of the features of the arguments. Features are selected so they correspond to natural parameters that can take on any real number.

Let $\psi(x)$ be the features of x . Then

$$\eta(x) = \mathbf{N}^T \begin{bmatrix} 1 \\ \psi(x) \end{bmatrix}$$

$$p_{\mathbf{N}}(y|x) = \exp\left(\begin{bmatrix} 1 \\ \psi(x) \end{bmatrix}^T \mathbf{N} \phi(y) - g\left(\mathbf{N}^T \begin{bmatrix} 1 \\ \psi(x) \end{bmatrix}\right)\right)$$

where \mathbf{N} is a matrix containing our parameters. It must have 0 for each entry whose row corresponds to a sufficient statistic of y that is not a feature and whose column is not 1. This ensures that only the natural parameters that are features of y are affected by x .

Additionally, we provide a `randomValue` function for convenience, which returns a sample from the “default” distribution for a given type. Table 1 shows the types in Quipp and their corresponding sufficient statistics, features, and English descriptions.

Note that `Vector(n, T)` is shorthand for `Tuple([T, T, ..., T])` (with n copies of T), and `Bool` is shorthand for `Categorical(2)`.

4 Example: improving the clustering model

Using this knowledge about the form of the random functions, we can improve the clustering model from before. As observed in Figure 1, the two clusters are forced to have the same variance. They do not fit the data well, since the data has a different shape in each location. To fix this problem, we can substitute the following model, which uses a separate random function (and therefore a separate axis-aligned Gaussian distribution) for each cluster:

```

162 var Cluster = Categorical(2);
163 var Point = Vector(2, Double);
164 var getPointFunctions = [randFunction(Point), randFunction(Point)];
165
166 var model = function() {
167   var cluster = randomValue(Cluster);
168   observe(getPointFunctions[cluster]);
169 };

```




Figure 2: Clustering the Iris dataset with independent variances

Using this model, we get the clusters shown in Figure 2. This model (with the parameters filled in) is estimated to assign probability density e^{-337} to the data, yielding a perplexity of 0.1058. This is a significantly improved fit.

The fact that the gaussian distribution for each cluster must be axis-aligned limits the degree to which the model can fit the data. To allow each cluster to be an arbitrary Gaussian distributions (where the x and y coordinates may be correlated), we can use the following model:

```

199 var Cluster = Categorical(2);
200 var getX = [randFunction(Double), randFunction(Double)];
201 var getY = [randFunction(Double, Double), randFunction(Double, Double)];
202
203 var model = function() {
204   var cluster = randomInteger(nclusters);
205   var x = observe(getX[cluster]);
206   var y = observe(getY[cluster], x);
207 };

```

Here, there are 2 random functions for each cluster, one to get the x coordinate and one to get the y coordinate (whose distribution may depend linearly on the x coordinate). This allows x and y to be correlated, improving fit, as seen in Figure . This model is estimated to assign probability e^{-269} to the data, yielding a perplexity of 0.166. This is a large improvement from the previous model.

5 Inference

To infer both latent variables and parameters, we use a Monte Carlo expectation maximization algorithm [1] on the probabilistic model, iterating stages of estimating latent variables using Metropolis




Figure 3: Clustering the Iris dataset with multivariate Gaussian distributions

Hastings and inferring parameters using gradient descent. The first iteration uses randomly generated parameters.

For the expectation step, we must estimate latent variable distributions given fixed values for the parameters. To do this, we can replace unknown random functions in the Quipp program with random functions set to use these fixed parameter values, yielding a probabilistic program. We use the Metropolis Hastings algorithm to perform inference in this program, yielding a distribution of execution traces (where each execution trace specifies the result of every call to a random function). Next, for each random function, we can find all calls to it in the trace to get the training data.

For the maximization step, given samples from each random function, we set the parameters of the function to maximize the likelihood of the samples. To do this we, we use gradient descent.

Given (x, y) samples, parameter estimation to maximize log probability is a convex problem because the log probability function is concave as a function of \mathbf{N} :

$$\log p_{\mathbf{N}}(y|x) = \left[\begin{matrix} 1 \\ \psi(x) \end{matrix} \right]^T \mathbf{N} \phi(y) - g \left(\mathbf{N}^T \left[\begin{matrix} 1 \\ \psi(x) \end{matrix} \right] \right)$$

This relies on the fact that g is convex, but this is true in general for any exponential family distribution. Since the problem is convex, it is possible to use gradient descent to optimize the parameters. Although the only exponential family distributions we use in this paper are the categorical and Gaussian distributions, we can use the same algorithms for other exponential families, such as the Poisson and gamma distributions.

6 Evaluation

To evaluate performance, for each model, we:

- Randomly generate parameters θ
- Generate datasets x_{train}, x_{test} using θ
- Estimate $\log P(x_{test}|\theta)$
- Use the EM algorithm to infer approximate parameters $\hat{\theta}_i$ from x_{train} , for iteration $i = 1, 2, \dots$

- Estimate $\log P(x_{test}|\hat{\theta}_i)$
- Use the previous estimates to estimate the regret $\log P(x_{test}|\hat{\theta}_i) - \log P(x_{test}|\theta)$

Estimating $\log P(x_{test}|\theta)$ is nontrivial, given that the model contains latent variables. We use the Sequential Monte Carlo algorithm for this, as described in [2]. For the following examples, for each iteration except the first, we plot quantiles of average regret per sample over multiple runs. We exclude the first iteration because it represents the regret of randomly generated parameters, which is much higher than the regret in future iterations. Since the expectation maximization algorithm does not always converge to a global optimum, we do not expect regret to go to 0 as we increase the number of iterations.

7 Examples

7.1 Clustering

```

var Cluster = Categorical(3);
var Point = Vector(2, Double);
var getPoint = randFunction(Cluster, Point);

var model = function() {
  var cluster = randomValue(Cluster);
  observe(getPoint, cluster);
};

```

In this example, we cluster 2d points into 3 different clusters. Given a cluster, the distribution for a point is some independent Gaussian distribution. This is similar to fuzzy c-means clustering. Regret usually reduces quite a lot as iterations increase.




Figure 4: Regret-per-sample quantiles for clustering

7.2 Naive Bayes

```

var Class = Categorical(2);
var Features = Vector(10, Bool);
var classFeatures = [randFunction(Features), randFunction(Features)];

var model = function() {

```

```

324     var whichClass = observe(randomValue, Class);
325     observe(classFeatures[whichClass]);
326 };
327

```

The naive Bayes model is similar to the clustering model. We have two classes and a feature distribution for each. Since each feature is boolean, we will learn a different categorical distribution for each class. The main difference from clustering is that, since this is a supervised learning method, we observe the class in addition to the features. Since there are no latent variables, the system finds the correct parameters in a single iteration, so the regret plot is uninteresting.

7.3 Factor analysis

```

335     var Factors = Vector(2, Double);
336     var Point = Vector(5, Double);
337     var getPoint = randFunction(Factors, Point);
338
339     var model = function() {
340         var factors = randomValue(Factors);
341         return observe(getPoint, factors);
342     };
343

```

The factor analysis model is very similar to the clustering model. The main difference is that we replace the categorical `ClusterType` type with a vector type. `randomValue(Factors)` will return a vector from the standard multivariate normal distribution. This results in the model attempting to predict each point as an affine function of a vector of standard normal values, with Gaussian noise. Regret reduces as iterations increase.




Figure 5: Regret-per-sample quantiles for factor analysis

7.4 Hidden Markov model

```

372     var chainLength = 20;
373     var State = Categorical(2);
374     var Obs = Categorical(4);
375     var transFun = randFunction(State, State);
376     var obsFun = randFunction(State, Obs);
377
378     var observeStates = function(startState, i) {

```

```

378     if (i == chainLength) {
379         return [];
380     } else {
381         observe(obsFun, startState);
382         observeStates(transFun(startState), i+1);
383     }
384 };
385
386 var model = function() {
387     return observeStates(randomValue(State), 0);
388 };

```

In this example, we use the unknown function `transFun` for state transitions and `obsFun` for observations. This means that we will learn both the state transition matrix and the observation matrix. Regret starts low and does not change over time, indicating that only one EM iteration was necessary to accurately estimate the parameters.




Figure 6: Regret-per-sample quantiles for hidden Markov model

7.5 Latent Dirichlet allocation

```

418 var maxWordsPerDocument = 100;
419 var Topic = Categorical(3);
420 var Word = Categorical(10);
421 var topicToWord = randFunction(Topic, Word);
422
423 var model = function() {
424     var whichClass = randomValue(Topic);
425     var nWords = observe(randomIntegerERP, maxWordsPerDocument);
426     repeat(nWords, function(wordIndex) {
427         observe(classToWord, whichClass);
428     });
429 };

```

For latent Dirichlet allocation, we use the unknown function `topicToWord` to map latent topics to word distributions. We will learn a different categorical distribution for each topic. Regret generally reduces over iterations.




Figure 7: Regret-per-sample quantiles for latent Dirichlet allocation

8 Discussion

It is possible to write many useful machine learning models as Quipp programs and then use a generic Monte Carlo expectation maximization algorithm for inference. This algorithm infers parameters reasonably well. Non-experts will find it easier to write useful machine learning models in this language compared to other probabilistic programming languages.

In the future, it will be useful to support additional types. For example, we could add types for standard exponential families such as Poisson and Gamma. Supporting disjoint union types is also straightforward. Also, it will be useful to create a more user-friendly interface for inferring parameters and performing additional data processing, such as inferring latent values given fixed parameters or visualizing the distributions implied by the parameters.

References

- [1] Christophe Andrieu, Nando de Freitas, Arnaud Doucet, and Michael I. Jordan. An introduction to mcmc for machine learning. *Machine Learning*, 50(1-2):5–43, 2003.
- [2] Noah D Goodman and Andreas Stuhlmüller. The Design and Implementation of Probabilistic Programming Languages. <http://dippl.org>, 2014. Accessed: 2015-6-4.
- [3] M. Lichman. UCI machine learning repository, 2013.